

• 1. GROUP MEMBERS(Name AND Roll number)

Ashutosh(20/49052)

Monu Kumar(20/49022)

Shambhu(20/49073)

• 2. MOBILE PRICE CLASSIFICATION

SOURCE:-<https://www.kaggle.com/datasets/iabhishekofficial/mobile-price-classification?select=test.csv>

DESCRIPTION:- IN the given dataset **Mobile Price Classification** There are 20 attributes(Columns) and 1 classification attributes having 4 classes((low cost=0), (medium cost=1), (high cost=2) and (very high cost=4)) and 2000 records.

According to the attributes we predict the class of the mobile phone.

Attributes(Columns):-

1. **battery_power**(Total energy a battery can store in one time measured in mAh).
2. **blue**(Has bluetooth or not).
3. **clock_speed**(speed at which microprocessor executes instructions)
4. **dual_sim**(Has dual sim support or not).
5. **fc**(Front Camera mega pixels).
6. **four_g**(Has 4G or not).
7. **int_memory**(Internal Memory in Gigabytes)
8. **m_dep**(Mobile Depth in cm).
9. **mobile_wt**(Weight of mobile phone).
10. **n_cores**(Number of cores of processor).
11. **pc**(Primary Camera mega pixels).
12. **px_height**(Pixel Resolution Height).
13. **px_width**(Pixel Resolution Width).
14. **ram**(Random Access Memory in Mega Bytes)
15. **sc_h**(Screen Height of mobile in cm).
16. **sc_w**(Screen Width of mobile in cm).
17. **talk_time**(longest time that a single battery charge will last when you are).
18. **three_g**(Has 3G or not).
19. **touch_screen**(Has touch screen or not).
20. **wifi**(Has wifi or not).
21. **price_range**(This is the target variable with value of 0(low cost), 1(medium cost), 2(high cost) and 3(very high cost).)

```
# LIBRARIES
# !pip install fabulous
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots
from fabulous import text as ft
```

```
import pprint as pprint
import seaborn as sns
from fabulous.color import *
```

```
#READ CSV FILE
df_Original=pd.read_csv("train.csv")
print(bold(yellow(df_Original)))
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	\
0	842	0	2.2	0	1	0	7	
1	1021	1	0.5	1	0	1	53	
2	563	1	0.5	1	2	1	41	
3	615	1	2.5	0	0	0	10	
4	1821	1	1.2	0	13	1	44	
...	
1995	794	1	0.5	1	0	1	2	

```

1996    1965    1      2.6    1  0    0    39
1997    1911    0      0.9    1  1    1    36
1998    1512    0      0.9    0  4    1    46
1999    510     1      2.0    1  5    1    45

   m_dep  mobile_wt  n_cores ... px_height  px_width  ram  sc_h  sc_w \
0     0.6       188      2 ...        20       756  2549    9    7
1     0.7       136      3 ...       905      1988  2631   17    3
2     0.9       145      5 ...      1263      1716  2603   11    2
3     0.8       131      6 ...      1216      1786  2769   16    8
4     0.6       141      2 ...      1208      1212  1411    8    2
...
1995   0.8       106      6 ...      1222      1890   668   13    4
1996   0.2       187      4 ...       915      1965  2032   11   10
1997   0.7       108      8 ...       868      1632  3057   9    1
1998   0.1       145      5 ...       336      670   869   18   10
1999   0.9       168      6 ...      483      754  3919   19    4

  talk_time  three_g  touch_screen  wifi  price_range
0         19       0          0     1        1
1         7       1          1     0        2
2         9       1          1     0        2
3        11       1          0     0        2
4        15       1          1     0        1
...
1995   19       1          1     0        0
1996   16       1          1     1        2
1997   5        1          1     0        3
1998   19       1          1     1        0
1999   2        1          1     1        3

```

[2000 rows x 21 columns]

#Describe

```
print(bold(blue(df_Original.describe())))
```

	battery_power	blue	clock_speed	dual_sim	fc	\	
count	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000		
mean	1238.518500	0.4950	1.522250	0.509500	4.309500		
std	439.418206	0.5001	0.816004	0.500035	4.341444		
min	501.000000	0.0000	0.500000	0.000000	0.000000		
25%	851.750000	0.0000	0.700000	0.000000	1.000000		
50%	1226.000000	0.0000	1.500000	1.000000	3.000000		
75%	1615.250000	1.0000	2.200000	1.000000	7.000000		
max	1998.000000	1.0000	3.000000	1.000000	19.000000		
	four_g	int_memory	m_dep	mobile_wt	n_cores	...	\
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	...	
mean	0.521500	32.046500	0.501750	140.249000	4.520500	...	
std	0.499662	18.145715	0.288416	35.399655	2.287837	...	
min	0.000000	2.000000	0.100000	80.000000	1.000000	...	
25%	0.000000	16.000000	0.200000	109.000000	3.000000	...	
50%	1.000000	32.000000	0.500000	141.000000	4.000000	...	
75%	1.000000	48.000000	0.800000	170.000000	7.000000	...	
max	1.000000	64.000000	1.000000	200.000000	8.000000	...	
	px_height	px_width	ram	sc_h	sc_w	\	
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000		
mean	645.108000	1251.515500	2124.213000	12.306500	5.767000		
std	443.780811	432.199447	1084.732044	4.213245	4.356398		
min	0.000000	500.000000	256.000000	5.000000	0.000000		
25%	282.750000	874.750000	1207.500000	9.000000	2.000000		
50%	564.000000	1247.000000	2146.500000	12.000000	5.000000		
75%	947.250000	1633.000000	3064.500000	16.000000	9.000000		
max	1960.000000	1998.000000	3998.000000	19.000000	18.000000		
	talk_time	three_g	touch_screen	wifi	price_range		
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000		
mean	11.011000	0.761500	0.503000	0.507000	1.500000		
std	5.463955	0.426273	0.500116	0.500076	1.118314		
min	2.000000	0.000000	0.000000	0.000000	0.000000		
25%	6.000000	1.000000	0.000000	0.000000	0.750000		
50%	11.000000	1.000000	1.000000	1.000000	1.500000		
75%	16.000000	1.000000	1.000000	1.000000	2.250000		
max	20.000000	1.000000	1.000000	1.000000	3.000000		

[8 rows x 21 columns]

#INFO

```
print(bold(magenta(df_Original.info())))
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   battery_power  2000 non-null  int64  
 1   blue           2000 non-null  int64  
 2   clock_speed   2000 non-null  float64
 3   fc             2000 non-null  float64
 4   four_g         2000 non-null  float64
 5   int_memory     2000 non-null  float64
 6   m_dep          2000 non-null  float64
 7   mobile_wt      2000 non-null  float64
 8   n_cores         2000 non-null  float64
 9   px_height      2000 non-null  float64
 10  px_width       2000 non-null  float64
 11  ram            2000 non-null  float64
 12  sc_h           2000 non-null  float64
 13  sc_w           2000 non-null  float64
 14  talk_time      2000 non-null  float64
 15  three_g        2000 non-null  float64
 16  touch_screen   2000 non-null  float64
 17  wifi            2000 non-null  float64
 18  price_range    2000 non-null  float64
```

```

3 dual_sim      2000 non-null   int64
4 fc           2000 non-null   int64
5 four_g       2000 non-null   int64
6 int_memory    2000 non-null   int64
7 m_dep        2000 non-null   float64
8 mobile_wt     2000 non-null   int64
9 n_cores      2000 non-null   int64
10 pc           2000 non-null   int64
11 px_height    2000 non-null   int64
12 px_width     2000 non-null   int64
13 ram          2000 non-null   int64
14 sc_h         2000 non-null   int64
15 sc_w         2000 non-null   int64
16 talk_time    2000 non-null   int64
17 three_g      2000 non-null   int64
18 touch_screen 2000 non-null   int64
19 wifi          2000 non-null   int64
20 price_range   2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
None

```

```
#HEAD
print(bold(df_Original.head()))
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	\
0	842	0	2.2	0	1	0	7	0.6	
1	1021	1	0.5	1	0	1	53	0.7	
2	563	1	0.5	1	2	1	41	0.9	
3	615	1	2.5	0	0	0	10	0.8	
4	1821	1	1.2	0	13	1	44	0.6	

	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w	talk_time	\
0	188	2	...	20	756	2549	9	7	19	
1	136	3	...	905	1988	2631	17	3	7	
2	145	5	...	1263	1716	2603	11	2	9	
3	131	6	...	1216	1786	2769	16	8	11	
4	141	2	...	1208	1212	1411	8	2	15	

	three_g	touch_screen	wifi	price_range
0	0	0	1	1
1	1	1	0	2
2	1	1	0	2
3	1	0	0	2
4	1	1	0	1

```
[5 rows x 21 columns]
```

```
#TAIL
print(bold(df_Original.tail()))



|      | battery_power | blue | clock_speed | dual_sim | fc | four_g | int_memory | \ |
|------|---------------|------|-------------|----------|----|--------|------------|---|
| 1995 | 794           | 1    | 0.5         | 1        | 0  | 1      | 2          |   |
| 1996 | 1965          | 1    | 2.6         | 1        | 0  | 0      | 39         |   |
| 1997 | 1911          | 0    | 0.9         | 1        | 1  | 1      | 36         |   |
| 1998 | 1512          | 0    | 0.9         | 0        | 4  | 1      | 46         |   |
| 1999 | 510           | 1    | 2.0         | 1        | 5  | 1      | 45         |   |


|      | m_dep | mobile_wt | n_cores | ... | px_height | px_width | ram  | sc_h | sc_w | \ |
|------|-------|-----------|---------|-----|-----------|----------|------|------|------|---|
| 1995 | 0.8   | 106       | 6       | ... | 1222      | 1890     | 668  | 13   | 4    |   |
| 1996 | 0.2   | 187       | 4       | ... | 915       | 1965     | 2032 | 11   | 10   |   |
| 1997 | 0.7   | 108       | 8       | ... | 868       | 1632     | 3057 | 9    | 1    |   |
| 1998 | 0.1   | 145       | 5       | ... | 336       | 670      | 869  | 18   | 10   |   |
| 1999 | 0.9   | 168       | 6       | ... | 483       | 754      | 3919 | 19   | 4    |   |


|      | talk_time | three_g | touch_screen | wifi | price_range |
|------|-----------|---------|--------------|------|-------------|
| 1995 | 19        | 1       | 1            | 0    | 0           |
| 1996 | 16        | 1       | 1            | 1    | 2           |
| 1997 | 5         | 1       | 1            | 0    | 3           |
| 1998 | 19        | 1       | 1            | 1    | 0           |
| 1999 | 2         | 1       | 1            | 1    | 3           |


```

```
[5 rows x 21 columns]
```

```
#INDEX | COLUMNS
print(bold(red(df_Original.columns)))
print(blue(df_Original.index))
```

```
Index(['battery_power', 'blue', 'clock_speed', 'dual_sim', 'fc', 'four_g',
       'int_memory', 'm_dep', 'mobile_wt', 'n_cores', 'pc', 'px_height',
       'px_width', 'ram', 'sc_h', 'sc_w', 'talk_time', 'three_g',
       'touch_screen', 'wifi', 'price_range'],
      dtype='object')
RangeIndex(start=0, stop=2000, step=1)
```

```
#DIMENSION
print(bold(yellow("Dimension of Original data : ",np.ndim(df_Original))))
print(bold(red("Shape of Original data : ",np.shape(df_Original))))
print(bold(cyan("Type : ",type(df_Original))))
print(bold(magenta("DataTypes --->\n",df_Original.dtypes)))

Dimension of Original data : 2
Shape of Original data : (2000, 21)
Type : <class 'pandas.core.frame.DataFrame'>
DataTypes --->
battery_power      int64
blue               int64
clock_speed       float64
dual_sim           int64
fc                 int64
four_g             int64
int_memory         int64
m_dep              float64
mobile_wt          int64
n_cores            int64
pc                 int64
px_height          int64
px_width           int64
ram                int64
sc_h               int64
sc_w               int64
talk_time          int64
three_g            int64
touch_screen       int64
wifi               int64
price_range        int64
dtype: object
```

#Check NULL Values

```
print(bold(green(df_Original.isnull().sum())))
```

```
battery_power    0
blue             0
clock_speed     0
dual_sim         0
fc               0
four_g           0
int_memory       0
m_dep            0
mobile_wt        0
n_cores          0
pc               0
px_height        0
px_width         0
ram              0
sc_h              0
sc_w              0
talk_time        0
three_g           0
touch_screen     0
wifi              0
price_range      0
dtype: int64
```

3. Exploratory Data Analysis

1.1 Data Distribution frequency plots. (HISTOGRAM)

```
df_Original.hist(bins=30,figsize=(25, 15) )
```

```

array([[[<Axes: title={'center': 'battery_power'}>,
        <Axes: title={'center': 'blue'}>,
        <Axes: title={'center': 'clock_speed'}>,
        <Axes: title={'center': 'dual_sim'}>,
        <Axes: title={'center': 'fc'}>],
       [<Axes: title={'center': 'four_g'}>,
        <Axes: title={'center': 'int_memory'}>,
        <Axes: title={'center': 'm_dep'}>,
        <Axes: title={'center': 'mobile_wt'}>,
        <Axes: title={'center': 'n_cores'}>],
       [<Axes: title={'center': 'pc'}>,
        <Axes: title={'center': 'px_height'}>,
        <Axes: title={'center': 'px_width'}>,
        <Axes: title={'center': 'ram'}>,
        <Axes: title={'center': 'sc_h'}>],
       [<Axes: title={'center': 'sc_w'}>,
        <Axes: title={'center': 'talk_time'}>,
        <Axes: title={'center': 'three_g'}>,
        <Axes: title={'center': 'touch_screen'}>,
        <Axes: title={'center': 'wifi'}>],
       [<Axes: title={'center': 'price_range'}>, <Axes: >, <Axes: >,
        <Axes: >, <Axes: >]], dtype=object)

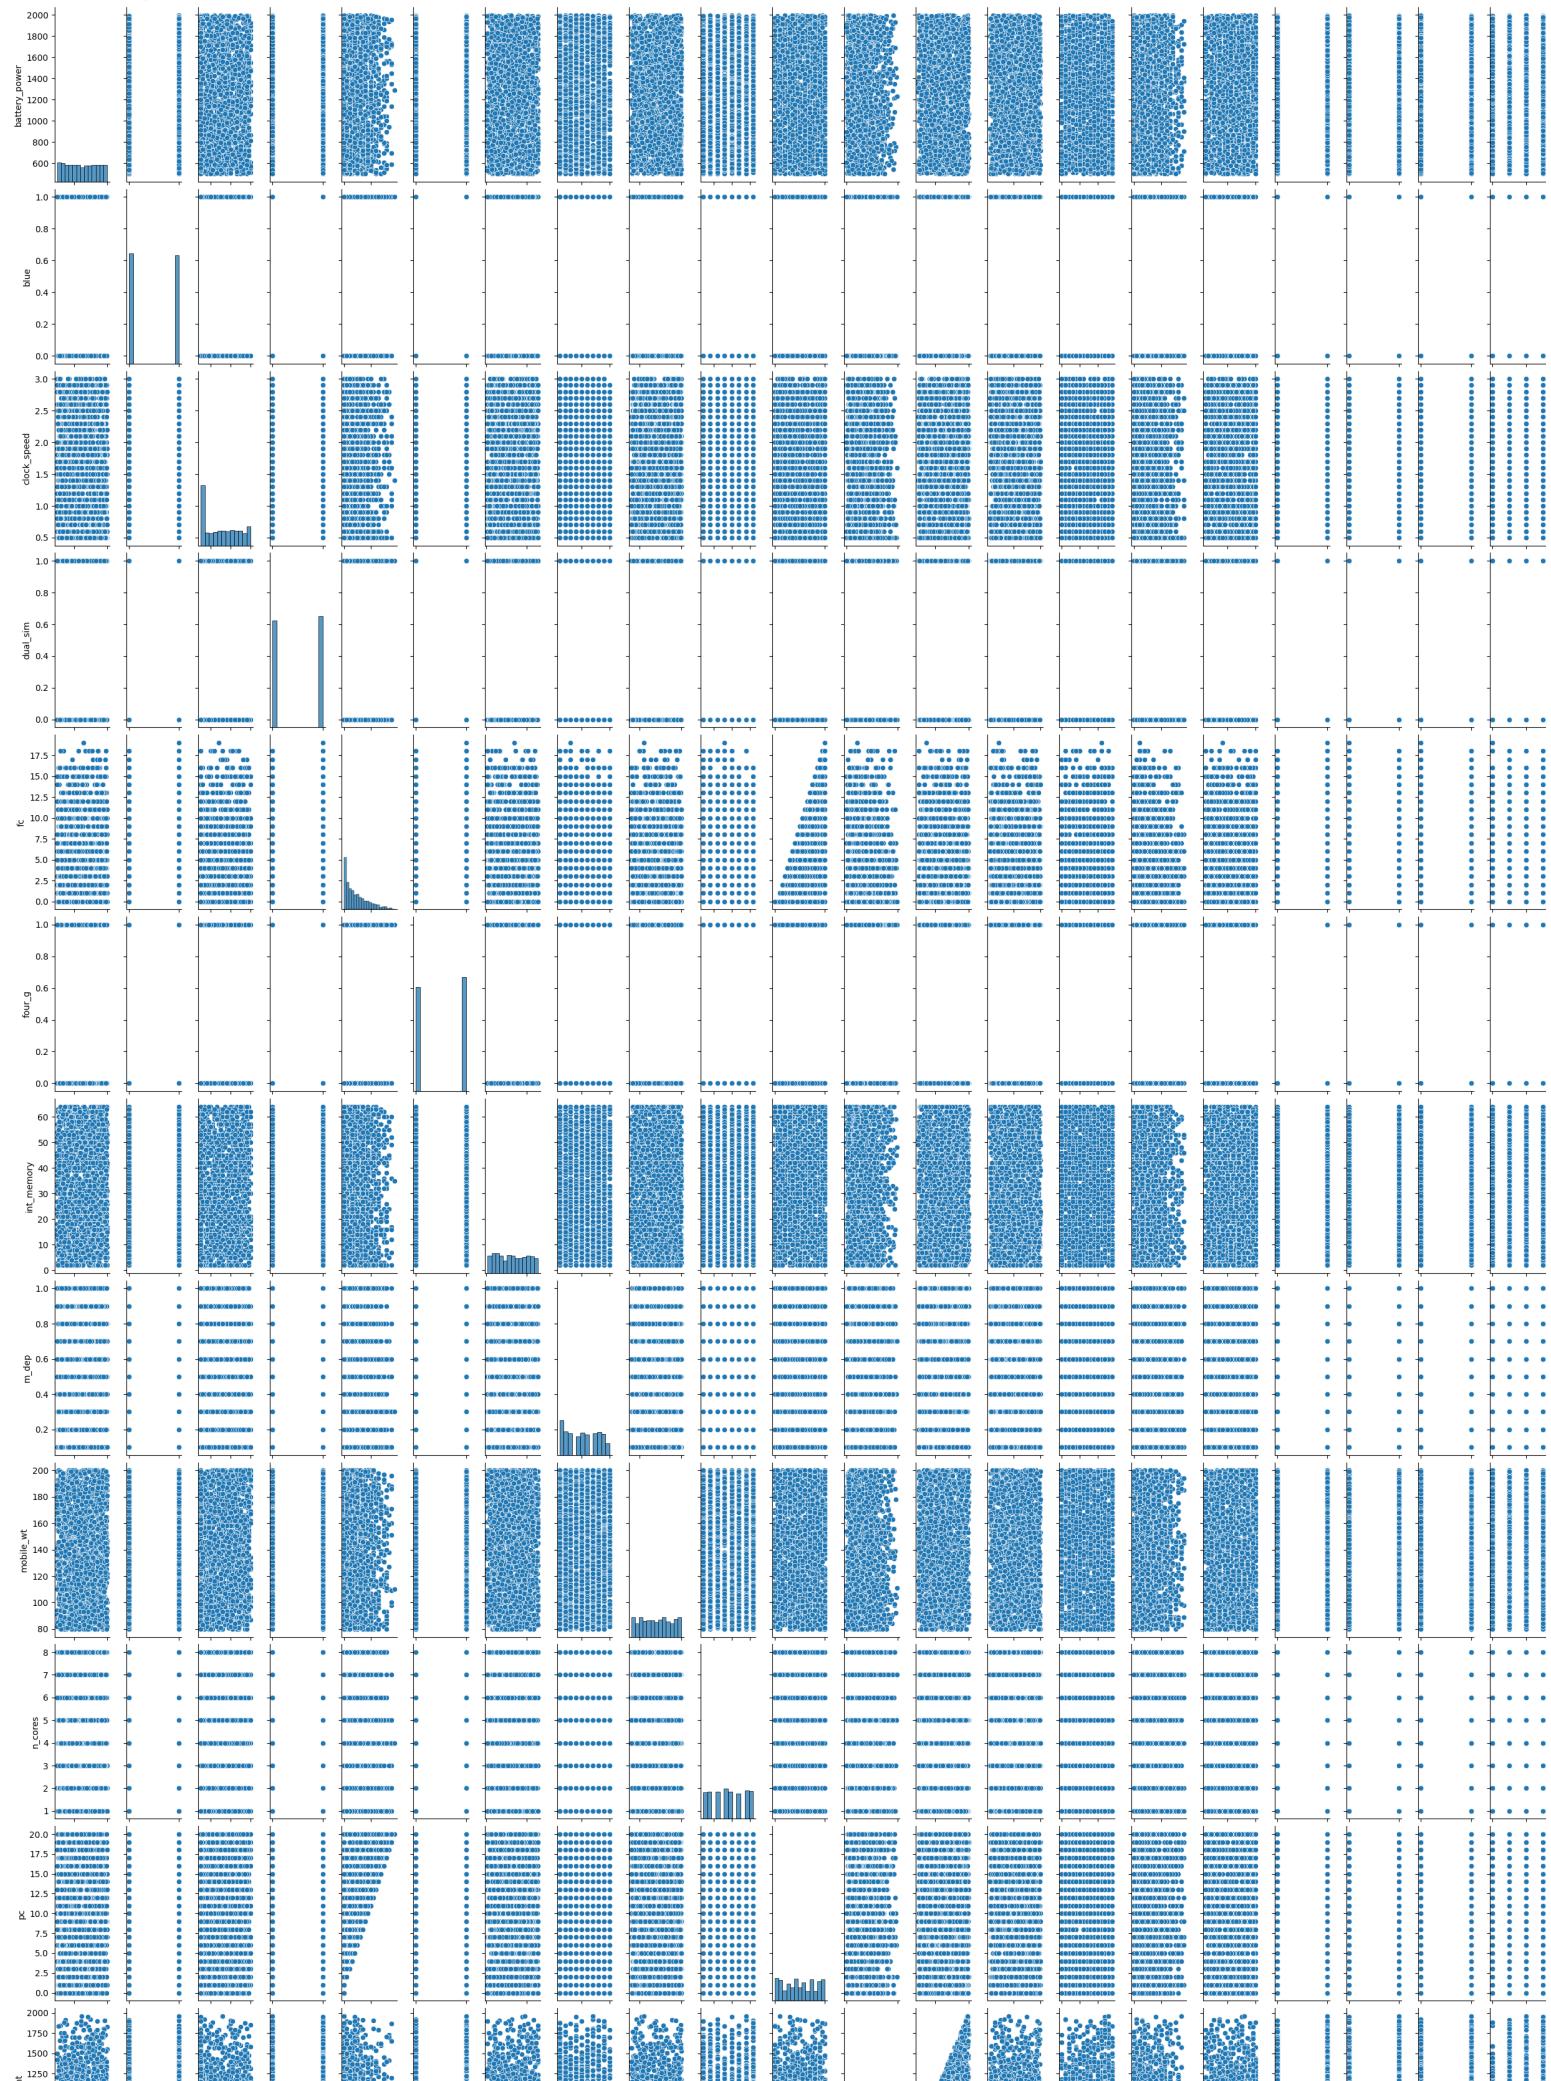
```



1.2 Scatter plot

```
sns.pairplot(df_Original,height=3,aspect=0.4,)
```

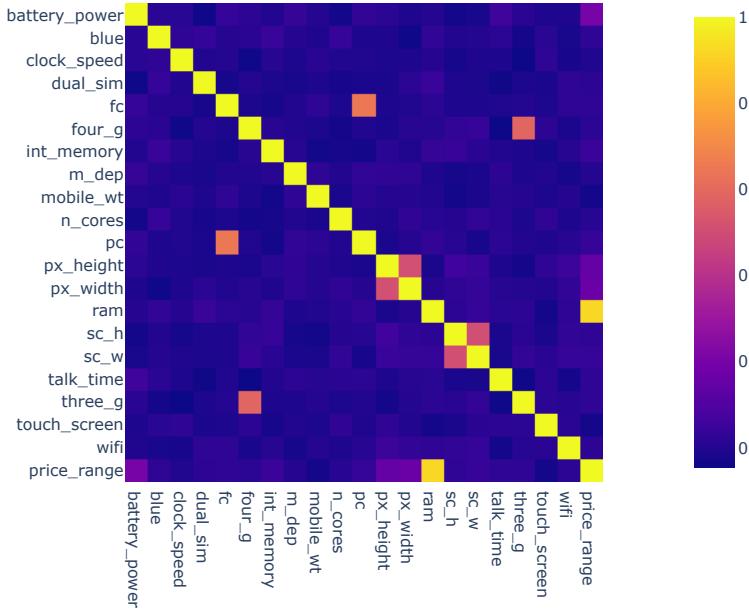
<seaborn.axisgrid.PairGrid at 0x7f1fcf8c6850>



1.3 Pearson correlation

```
# sns.heatmap(df_Original.corr(),cmap="RdYlBu")
```

```
fig = px.imshow(df_Original.corr())
fig.show()
```



```
print("Top 10 corelated Columns to the Tageted value : ")
```

```
abs(df_Original.corr()["price_range"]).nlargest(10)
```

```
Top 10 corelated Columns to the Tageted value :
```

```
price_range      1.000000
ram             0.917046
battery_power   0.200723
px_width        0.165818
px_height       0.148858
int_memory      0.044435
sc_w            0.038711
pc              0.033599
touch_screen    0.030411
mobile_wt        0.030302
Name: price_range, dtype: float64
```

```
#Attribute Classification
```

```
Independent_Attributes=df_Original[df_Original.corr()["price_range"].nlargest(10).index[1:]]
```

```
Dependent_Attribute=df_Original.iloc[:, -1]
```

```
print("Independent Attributes are :- /n")
```

```
print(Independent_Attributes.info())
```

```
Independent Attributes are :- /n
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   ram              2000 non-null   int64  
 1   battery_power    2000 non-null   int64  
 2   px_width         2000 non-null   int64  
 3   px_height        2000 non-null   int64  
 4   int_memory       2000 non-null   int64  
 5   sc_w             2000 non-null   int64  
 6   pc               2000 non-null   int64  
 7   three_g          2000 non-null   int64  
 8   sc_h             2000 non-null   int64  
dtypes: int64(9)
memory usage: 140.8 KB
None
```

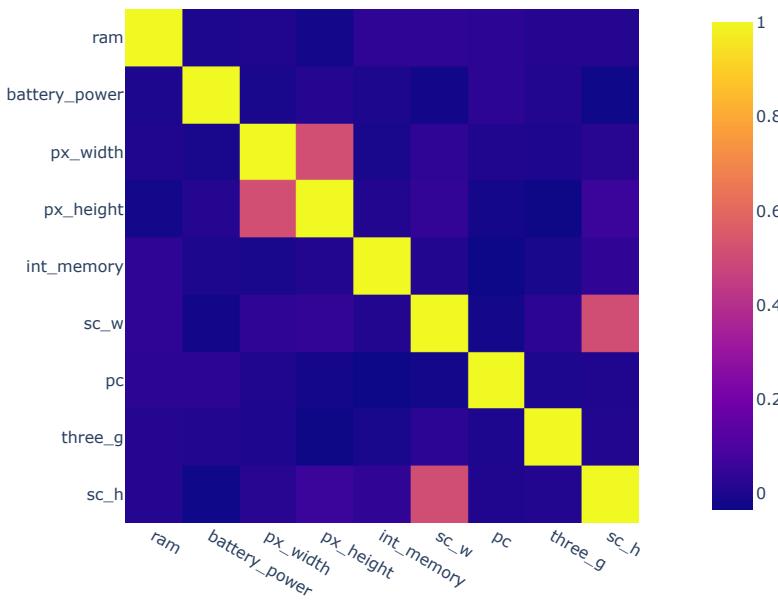
```
fig = px.scatter_matrix(Independent_Attributes, dimensions=(Independent_Attributes), height=1500, width=2500)
fig.update_traces(diagonal_visible=False)
fig.show()
```

⟳

```
4000
```

PEARSON CORELATION OF INDEPENDENT ATTRIBUTES:-

```
fig = px.imshow(Independent_Attributes.corr())
fig.show()
```



```
print("Dependent Attribute is :- \n")
print(Dependent_Attribute.info())
```

Dependent Attribute is :-

```
<class 'pandas.core.series.Series'>
RangeIndex: 2000 entries, 0 to 1999
Series name: price_range
Non-Null Count Dtype
-----
2000 non-null    int64
dtypes: int64(1)
memory usage: 15.8 KB
None
```

Dependent_Attribute

```
0      1
1      2
2      2
3      2
4      1
..
1995   0
1996   2
1997   3
1998   0
1999   3
Name: price_range, Length: 2000, dtype: int64
```

2.1 Removing Duplicate values

```
Duplicates_value=df_Original[Independent_Attributes.duplicated()]
Duplicates_value
```

```
battery_power blue clock_speed dual_sim fc four_g int_memory m_dep mobile_wt n_cores px_height px_width ram sc_h sc_w talk_time
```

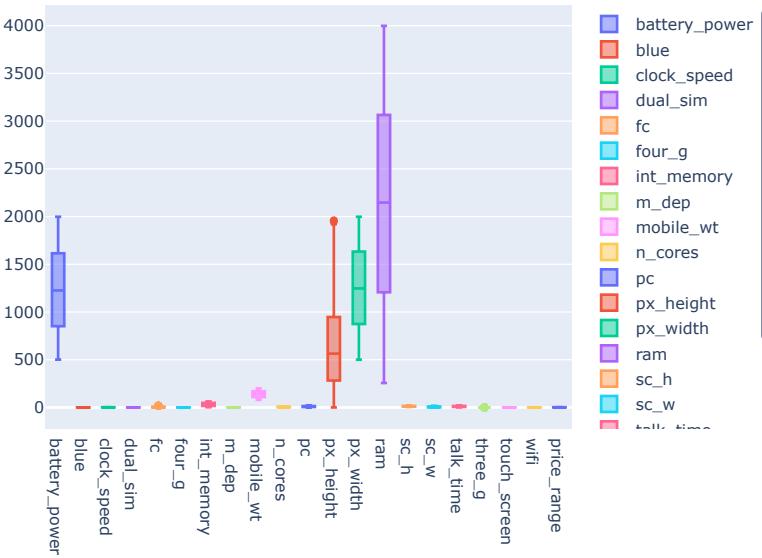
2.2 Handling Outliers

“”

```
fig = go.Figure()

for col in df_Original:
    fig.add_trace(go.Box(y=df_Original[col].values, name=df_Original[col].name))

fig.show()
```



2.3 Handling Null values (K-nearest neighbour method)

```
print(df_Original.isna().sum(), bold(yellow("\n\n\n NO NULL VALUE PRESENT :-")))
```

```
battery_power      0
blue                0
clock_speed         0
dual_sim            0
fc                  0
four_g              0
int_memory          0
m_dep               0
mobile_wt            0
n_cores             0
pc                  0
px_height           0
px_width            0
ram                 0
sc_h                0
sc_w                0
talk_time           0
three_g             0
touch_screen        0
wifi                0
price_range          0
dtype: int64
```

NO NULL VALUE PRESENT :-

3. Feature Scaling

3.1 STANDARD SCALER

```
from sklearn.preprocessing import StandardScaler
def normalize(Independent_Attributes):
    print("Mean and Standard Deviation Before")
    print(Independent_Attributes.mean(axis=0), Independent_Attributes.std(axis=0))
```

```

sc=StandardScaler()
Xscaled = sc.fit_transform(Independent_Attributes)

print("Mean and Standard Deviation After")
print(Xscaled.mean(axis=0).round(4), Xscaled.std(axis=0))
return Xscaled

```

```
# Standard=normalize(Independent_Attributes)
```

• Logistic Regression

```

from sklearn.metrics import confusion_matrix
from mlxtend.plotting import plot_confusion_matrix
from sklearn.metrics import classification_report
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.model_selection import StratifiedKFold
from sklearn .metrics import roc_auc_score
def Logistic_Regression(Independent_Attributes,Dependent_Attribute):
    X=Independent_Attributes.values
    Y=Dependent_Attribute.values
    print("*****Normalization/Standardization*****")
    Xscaled = normalize(X)

    skf=StratifiedKFold(n_splits=5, shuffle=True, random_state=2)
    acc=[]

    y_ori = np.array([], dtype=int)

    y_pre= np.array([], dtype=int)
    net_mat=np.zeros((4, 4))
    roc=[]
    for train_index, test_index in skf.split(X,Y):

        X_train=X[train_index]
        X_test=X[test_index]
        Y_train=Y[train_index]
        Y_test=Y[test_index]
        print(bold(green("Shape of X_train:- {} || X_test:- {} || Y_train:- {} || Y_test:- {}".format(X_train.shape,X_test.shape,Y_train.shape,Y_test.shape))))
        print(bold(magenta("Train:- {} || Test:- {}".format(np.bincount(Y[train_index]),np.bincount(Y[test_index])))))

        LRModel = LogisticRegression(max_iter=100000)
        LRModel.fit(X_train, Y_train)
        Y_testPred = LRModel.predict(X_test)

        y_ori=np.hstack((y_ori,Y_test))

        y_pre=np.hstack((y_pre,Y_testPred))
        testAccuracy = metrics.accuracy_score(Y_test, Y_testPred)
        print(bold(red("Test Accuracy", testAccuracy*100)))
        acc.append(testAccuracy)
        roc_score=roc_auc_score(Y_test, LRModel.predict_proba(X_test), multi_class='ovr')

        print(bold(yellow("ROC_AUC_SCORE: ",roc_score)))

        matrix1= confusion_matrix(Y_test, Y_testPred)
        #      sum of the total confusion matirx
        net_mat=net_mat+matrix1

        plot_confusion_matrix(matrix1,show_normed=True, colorbar=True, show_absolute=True)

        plt.show()
    print(bold(green(" Result Of Logsitic Regression: \n")))
    avg_accuracy=(sum(acc) / len(acc))*100
    print(bold(yellow_bg("average accuracy: ",avg_accuracy )))

    net_mat = net_mat.astype('int')
    print(bold(cyan(net_mat)))
    plot_confusion_matrix(net_mat,show_normed=True, colorbar=True, show_absolute=True, cmap='Blues')

    plt.show()

    roc_score=roc_auc_score(Y, LRModel.predict_proba(X), multi_class='ovr')

```

```
print(bold(yellow_bg("ROC_AUC_SCORE: ",roc_score)))\n\nprint("Classification Report:\n")\nreport=classification_report(y_ori, y_pre,output_dict=True)\nreport_Df=pd.DataFrame(report)\nprint(report_Df)\nsns.heatmap(report_Df.T,annot=True)\nreturn avg_accuracy\navg_accuracy_logistic=Logistic_Regression(Independent_Attributes,Dependent_Attribute)
```

```
*****Normalization/Standardization*****
Mean and Standard Deviation Before
[2.1242130e+03 1.2385185e+03 1.2515155e+03 6.4510800e+02 3.2046500e+01
 5.7670000e+00 9.9165000e+00 7.6150000e-01 1.2306500e+01] [1.08446083e+03 4.39308338e+02 4.32091384e+02
 1.81411780e+01 4.35530837e+00 6.06279867e+00 4.26166341e-01
 4.21219156e+00]
Mean and Standard Deviation After
[-0.  0.  0.  0. -0.  0.  0.  0.] [1.  1.  1.  1.  1.  1.  1.  1.]
Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:- (400,
Train:- [400 400 400 400] || Test:- [100 100 100 100]
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:
```

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of f AND g EVALUATIONS EXCEEDS LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

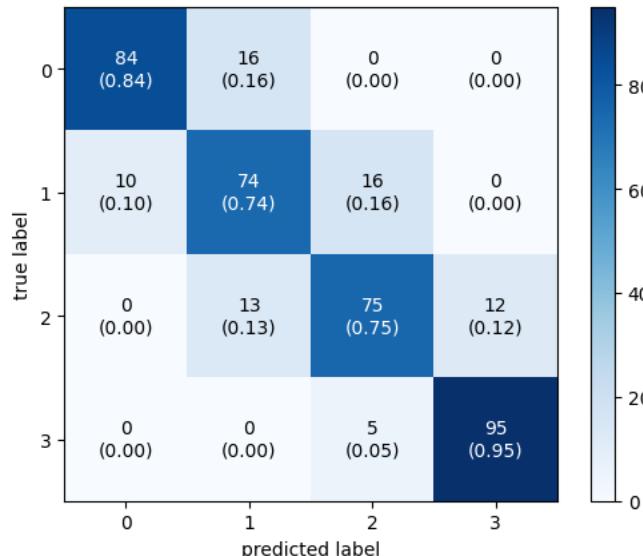
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

Test Accuracy82.0

ROC_AUC_SCORE: 0.9589083333333334



```
Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:- (400,
Train:- [400 400 400 400] || Test:- [100 100 100 100]
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:
```

lbfgs failed to converge (status=1):
STOP: TOTAL NO. of f AND g EVALUATIONS EXCEEDS LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

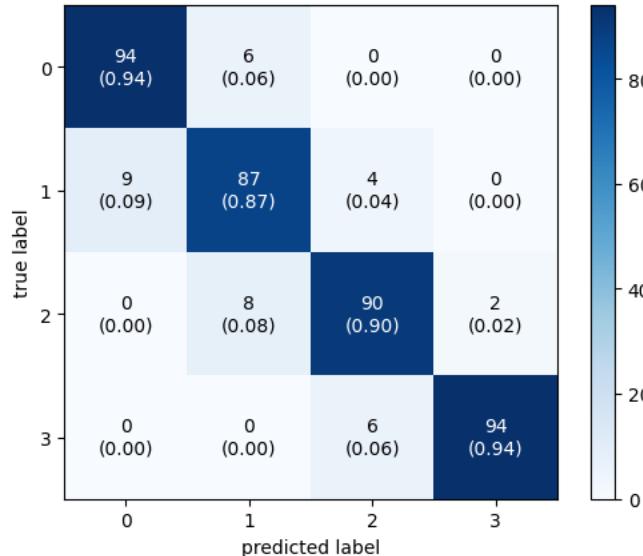
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

Test Accuracy91.25

ROC_AUC_SCORE: 0.9921333333333333



```
Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:- (400,
Train:- [400 400 400 400] || Test:- [100 100 100 100]
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:
```

lbfgs failed to converge (status=1):

▼ Descision Tree Classifier

Please also refer to the documentation for alternative solver options:

```
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.tree import export_graphviz
from six import StringIO
from IPython.display import Image
import pydotplus
def Decision_Tree(Independent_Attributes,Dependent_Attribute):
    X=Independent_Attributes.values
    Y=Dependent_Attribute.values
    print("*****Normalization/Standardization*****")
    XScaled = normalize(X)

    skf=StratifiedKFold(n_splits=5, shuffle=True, random_state=2)
    acc=[]

    y_ori = np.array([], dtype=int)

    y_pre= np.array([], dtype=int)
    net_mat=np.zeros((4, 4))

    for train_index, test_index in skf.split(X,Y):

        X_train=X[train_index]
        X_test=X[test_index]
        Y_train=Y[train_index]
        Y_test=Y[test_index]
        print(bold(green("Shape of X_train:- {} || X_test:- {} || Y_train:- {} || Y_test:- {}".format(X_train.shape,X_test.shape,Y_train.shape,Y_test.shape))))
        print(bold(magenta("Train:- {} || Test:- {}".format(np.bincount(Y[train_index]),np.bincount(Y[test_index])))))
        treemodel=DecisionTreeClassifier(criterion='entropy',max_depth=4)
        treemodel.fit(X_train, Y_train)

        Y_testPred =treemodel.predict(X_test)

        y_ori=np.hstack((y_ori,Y_test))

        y_pre=np.hstack((y_pre,Y_testPred))
        testAccuracy = metrics.accuracy_score(Y_test, Y_testPred)
        print(bold(red("Test Accuracy", testAccuracy*100)))
        acc.append(testAccuracy)

        roc_score=roc_auc_score(Y_test, treemodel.predict_proba(X_test), multi_class='ovr')

        print(bold(yellow("ROC_AUC_SCORE: ",roc_score)))

        matrix1= confusion_matrix(Y_test, Y_testPred)
        # sum of the total confusion matirx
        net_mat=net_mat+matrix1

    plot_confusion_matrix(matrix1,show_normed=True, colorbar=True, show_absolute=True)
    plt.show()

    plt.figure(figsize=(15,10))
    tree.plot_tree(treemodel,feature_names=Independent_Attributes.columns,filled=True)

    print(bold(green(" Result Of Decision Tree: \n")))
    avg_accuracy=(sum(acc) / len(acc))*100
    print(bold(yellow_bg("average accuracy: ",avg_accuracy )))

    net_mat = net_mat.astype('int')
    print(bold(cyan(net_mat)))
    plot_confusion_matrix(net_mat,show_normed=True, colorbar=True, show_absolute=True, cmap='Blues')

    plt.show()

    roc_score=roc_auc_score(Y, treemodel.predict_proba(X), multi_class='ovr')
    print(bold(yellow_bg("ROC_AUC_SCORE: ",roc_score)))
```

```
print(bold("Classification Report:\n"))
report=classification_report(y_ori, y_pre,output_dict=True)
report_Df=pd.DataFrame(report)
print(bold(magenta(report_Df)))
sns.heatmap(report_Df.T,annot=True)
return avg_accuracy

avg_accuracy_DCT=Decision_Tree(Independent_Attributes,Dependent_Attribute)
```

*****Normalization/Standardization*****

Mean and Standard Deviation Before

```
[2.1242130e+03 1.2385185e+03 1.2515155e+03 6.4510800e+02 3.2046500e+01  
5.7670000e+00 9.9165000e+00 7.6150000e-01 1.2306500e+01] [1.08446083e+03 4.39308338e+02 4.3209  
1.81411780e+01 4.35530837e+00 6.06279867e+00 4.26166341e-01  
4.21219156e+00]
```

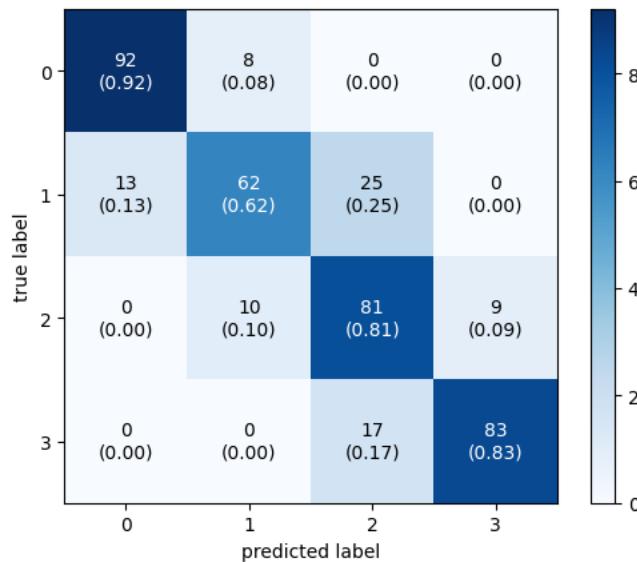
Mean and Standard Deviation After

```
[-0. 0. 0. 0. -0. -0. 0. 0.] [1. 1. 1. 1. 1. 1. 1. 1.]
```

Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:-
Train:- [400 400 400 400] || Test:- [100 100 100 100]

Test Accuracy 79.5

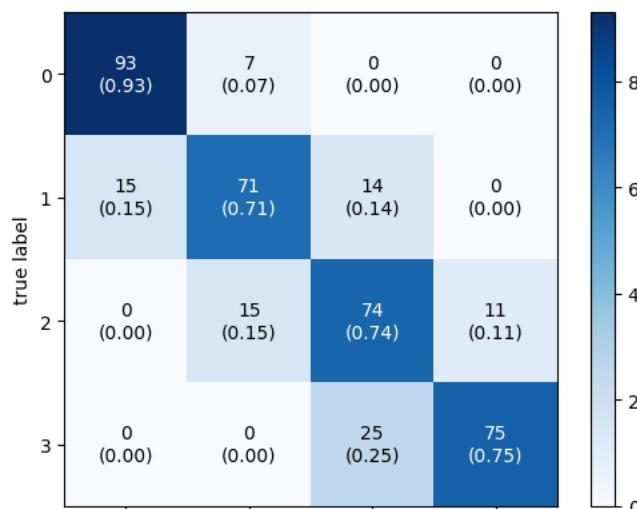
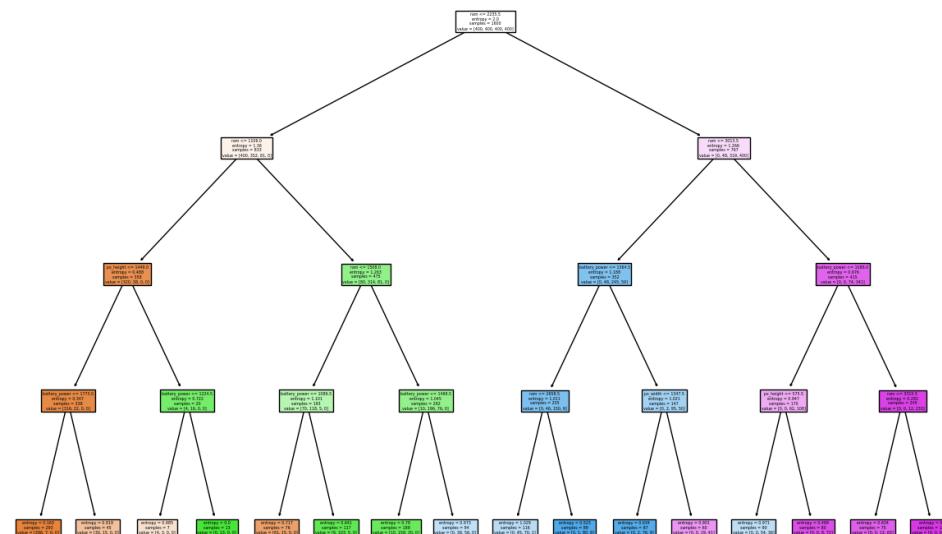
ROC_AUC_SCORE: 0.9398625



Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:-
Train:- [400 400 400 400] || Test:- [100 100 100 100]

Test Accuracy 78.25

ROC_AUC_SCORE: 0.9421708333333332



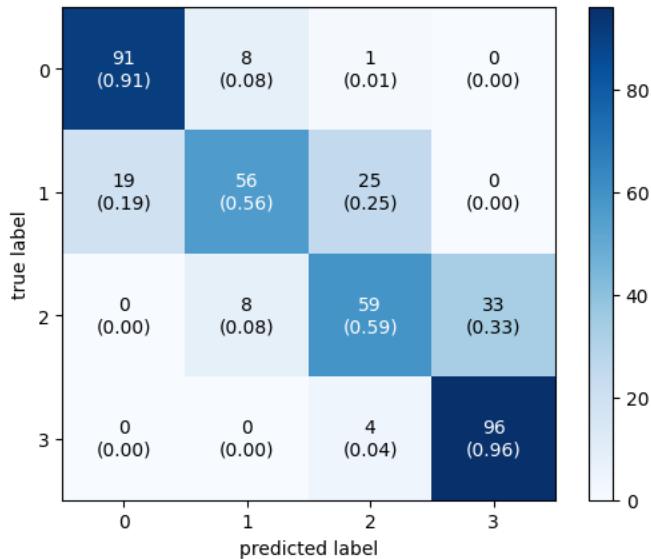
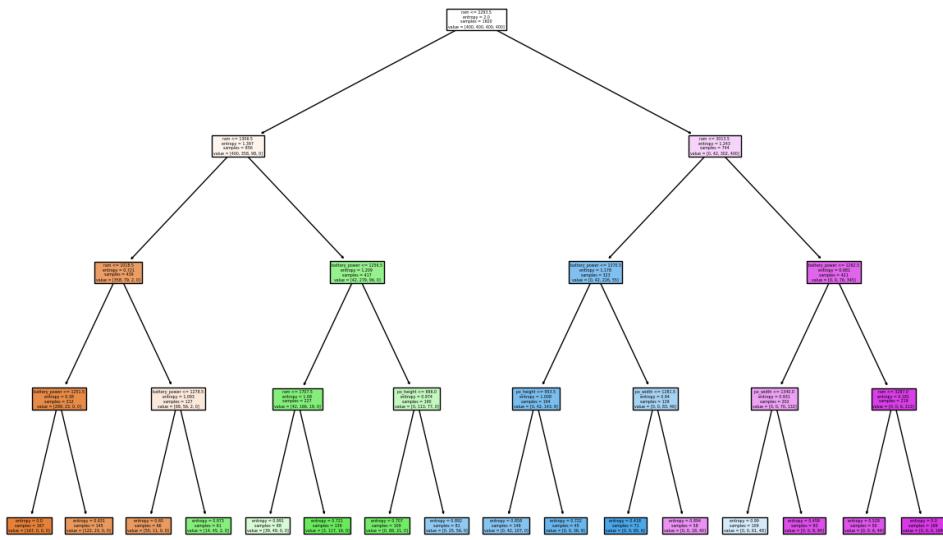
0 1 2 3
predicted label

Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:-

Train:- [400 400 400 400] || Test:- [100 100 100 100]

Test Accuracy75.5

ROC_AUC_SCORE: 0.9311666666666666

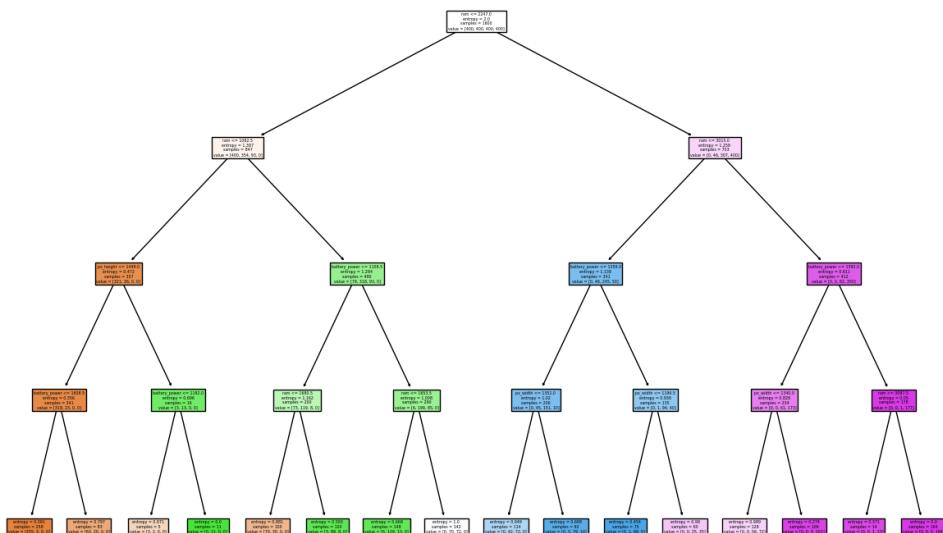


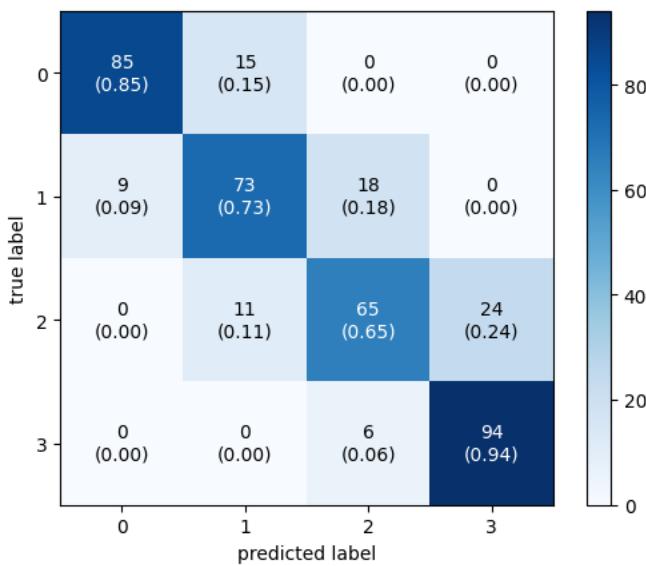
Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:-

Train:- [400 400 400 400] || Test:- [100 100 100 100]

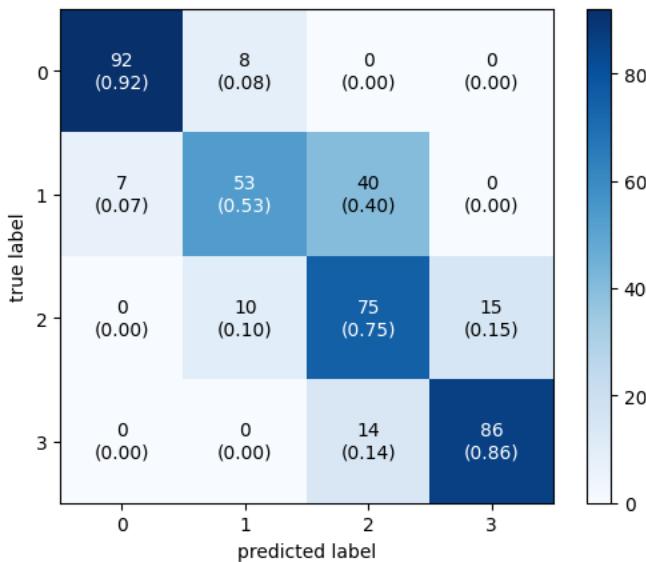
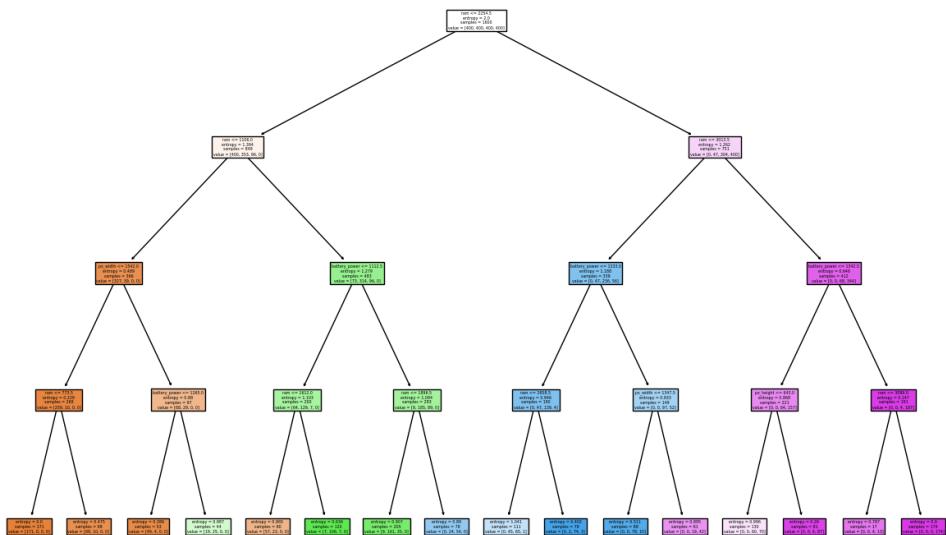
Test Accuracy79.25

ROC_AUC_SCORE: 0.9421375



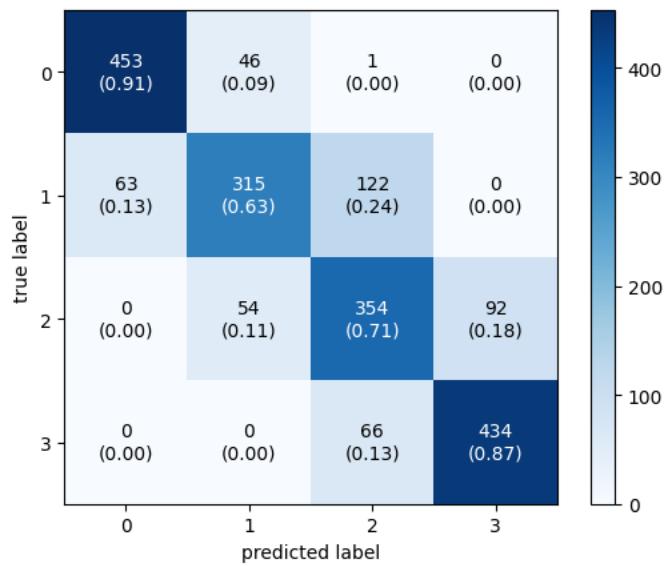
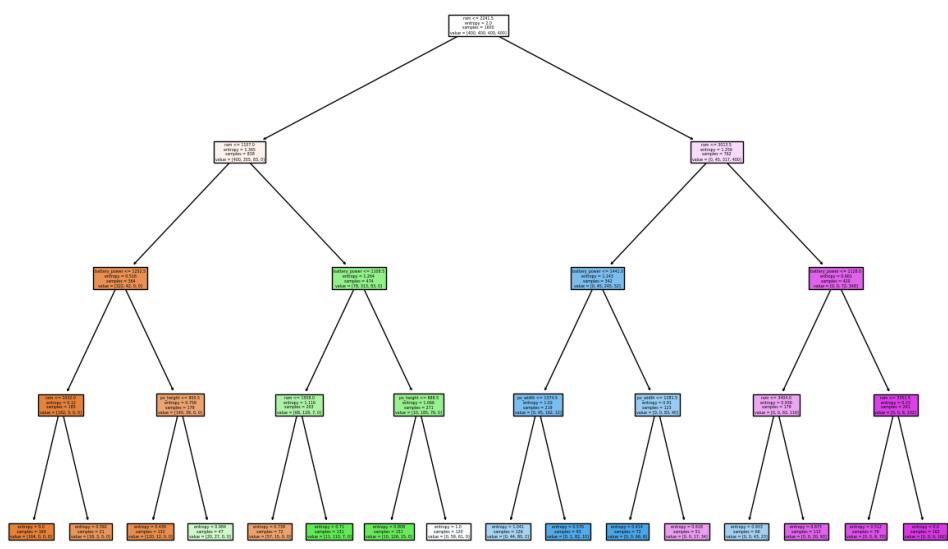


```
Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:-  
Train:- [400 400 400 400] || Test:- [100 100 100 100]  
Test Accuracy 76.5  
ROC_AUC_SCORE: 0.93815
```



Result Of Decision Tree:

```
average accuracy: 77.8
[[453 46 1 0]
 [ 63 315 122 0]
 [ 0 54 354 92]
 [ 0 0 66 434]]
```



ROC_AUC_SCORE: 0.955579

Classification Report:

	0	1	2	3	accuracy
precision	0.877907	0.759036	0.651934	0.825095	0.778
recall	0.906000	0.630000	0.708000	0.868000	0.778
f1-score	0.891732	0.688525	0.678811	0.846004	0.778
support	500.000000	500.000000	500.000000	500.000000	0.778

	macro avg	weighted avg
precision	0.778493	0.778493
recall	0.778000	0.778000
f1-score	0.776268	0.776268
support	2000.000000	2000.000000



- Gaussian Classification

```

from sklearn.naive_bayes import GaussianNB
def Gaussian_Naive_Bayes(Independent_Attributes,Dependent_Attribute):
    X=Independent_Attributes.values
    Y=Dependent_Attribute.values
    print("*****Normalization/Standardization*****")
    XScaled = normalize(X)

    skf=StratifiedKFold(n_splits=5, shuffle=True, random_state=2)
    acc=[]

    y_ori = np.array([], dtype=int)

    y_pre= np.array([], dtype=int)
    net_mat=np.zeros((4, 4))
    for train_index, test_index in skf.split(X,Y):

        X_train=X[train_index]
        X_test=X[test_index]
        Y_train=Y[train_index]
        Y_test=Y[test_index]
        print(bold(green("Shape of X_train:- {} || X_test:- {} || Y_train:- {} || Y_test:- {}".format(X_train.shape,X_test.shape,Y_train.shape,Y_test.shape))))
        print(bold(magenta("Train:- {} || Test:- {}".format(np.bincount(Y[train_index]),np.bincount(Y[test_index])))))

        GN = GaussianNB()
        GN.fit(X_train, Y_train)
        Y_testPred = GN.predict(X_test)

        y_ori=np.hstack((y_ori,Y_test))

        y_pre=np.hstack((y_pre,Y_testPred))

        testAccuracy = metrics.accuracy_score(Y_test, Y_testPred)
        print(bold(red("Test Accuracy", testAccuracy*100)))
        acc.append(testAccuracy)

        roc_score=roc_auc_score(Y_test, GN.predict_proba(X_test), multi_class='ovr')

        print(bold(yellow("ROC_AUC_SCORE: ",roc_score)))

        matrix1= confusion_matrix(Y_test, Y_testPred)
        # sum of the total confusion matirx
        net_mat=net_mat+matrix1

        plot_confusion_matrix(matrix1,show_normed=True, colorbar=True, show_absolute=True)

        plt.show()
        print(bold(green("Result Of Gaussian Naive Bayes Classifier: \n")))
        print(bold(magenta("Naive Bayes score: ",GN.score(X_test, Y_testPred))))
        avg_accuracy= (sum(acc) / len(acc))*100
        print(bold(cyan_bg("average accuracy: ", avg_accuracy)))

        net_mat = net_mat.astype('int')
        print(bold(cyan(net_mat)))
        plot_confusion_matrix(net_mat,show_normed=True, colorbar=True, show_absolute=True, cmap='Blues')

        plt.show()

        roc_score=roc_auc_score(Y,GN.predict_proba(X), multi_class='ovr')
        print(bold(yellow_bg("ROC_AUC_SCORE: ",roc_score)))

print("Classification Report:\n")
report=classification_report(y_ori, y_pre,output_dict=True)
report_Df=pd.DataFrame(report)
print(bold(blue(report_Df)))
sns.heatmap(report_Df.T,annot=True)
return avg_accuracy
avg_accuracy_GNB=Gaussian_Naive_Bayes(Independent_Attributes,Dependent_Attribute)

```

*****Normalization/Standardization*****

Mean and Standard Deviation Before

[2.1242130e+03 1.2385185e+03 1.2515155e+03 6.4510800e+02 3.2046500e+01
5.7670000e+00 9.9165000e+00 7.6150000e-01 1.2306500e+01] [1.08446083e+03 4.39308338e+02 4.320913e+01
1.81411780e+01 4.35530837e+00 6.06279867e+00 4.26166341e-01
4.21219156e+00]

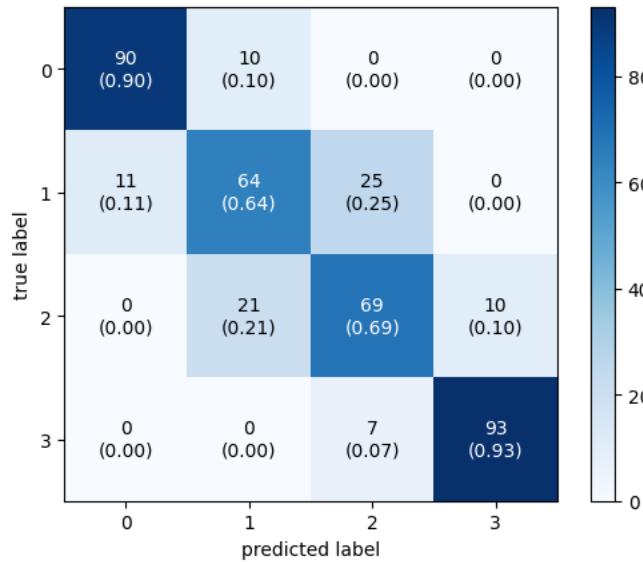
Mean and Standard Deviation After

[-0. 0. 0. 0. -0. -0. 0. 0.] [1. 1. 1. 1. 1. 1. 1. 1.]

Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:-
Train:- [400 400 400 400] || Test:- [100 100 100 100]

Test Accuracy79.0

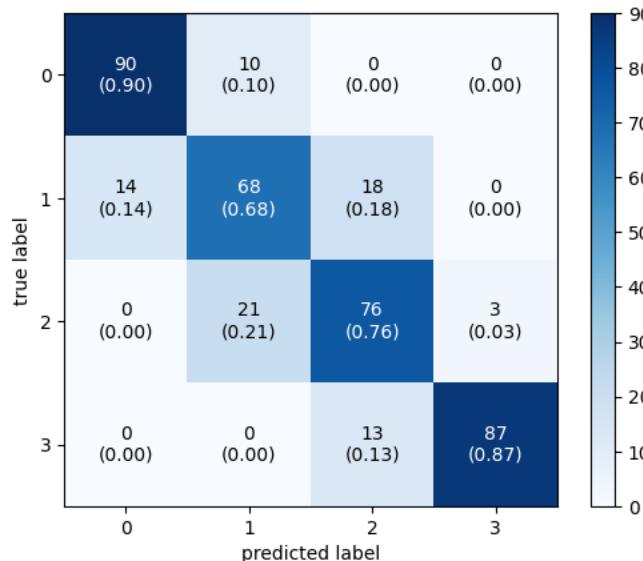
ROC_AUC_SCORE: 0.9438833333333333



Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:-
Train:- [400 400 400 400] || Test:- [100 100 100 100]

Test Accuracy80.25

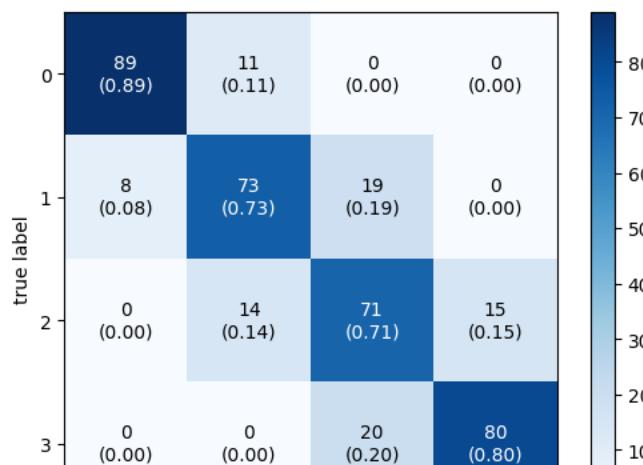
ROC_AUC_SCORE: 0.9558500000000001



Shape of X_train:- (1600, 9) || X_test:- (400, 9) || Y_train:- (1600,) || Y_test:-
Train:- [400 400 400 400] || Test:- [100 100 100 100]

Test Accuracy78.25

ROC_AUC_SCORE: 0.943



• KNN Classifier

```
from sklearn.neighbors import KNeighborsClassifier
def KNN(Independent_Attributes,Dependent_Attribute):
    X=Independent_Attributes.values
    Y=Dependent_Attribute.values
    print("*****Normalization/Standardization*****")
    XScaled = normalize(X)

    skf=StratifiedKFold(n_splits=5, shuffle=True, random_state=2)
    acc=[]

    y_ori = np.array([], dtype=int)

    y_pre= np.array([], dtype=int)
    net_mat=np.zeros((4, 4))
    for train_index, test_index in skf.split(X,Y):

        X_train=X[train_index]
        X_test=X[test_index]
        Y_train=Y[train_index]
        Y_test=Y[test_index]
        print(bold(green("Shape of X_train:- {} || X_test:- {} || Y_train:- {} || Y_test:- {}".format(X_train.shape,X_test.shape,Y_train.shape,Y_test.shape))))
        print(bold(magenta("Train:- {} || Test:- {}".format(np.bincount(Y[train_index]),np.bincount(Y[test_index])))))

        Kn = KNeighborsClassifier(n_neighbors=10)
        Kn.fit(X_train, Y_train)
        Y_testPred = Kn.predict(X_test)

        y_ori=np.hstack((y_ori,Y_test))

        y_pre=np.hstack((y_pre,Y_testPred))

        testAccuracy = metrics.accuracy_score(Y_test, Y_testPred)

        print(bold(red("Test Accuracy", testAccuracy*100)))
        acc.append(testAccuracy)

        roc_score=roc_auc_score(Y_test, Kn.predict_proba(X_test), multi_class='ovr')

        print(bold(yellow("ROC_AUC_SCORE: ",roc_score)))

        matrix1= confusion_matrix(Y_test, Y_testPred)
        # sum of the total confusion matirx
        net_mat=net_mat+matrix1

        plot_confusion_matrix(matrix1,show_normed=True, colorbar=True, show_absolute=True)

        plt.show()
        avg_accuracy=(sum(acc) / len(acc))*100
        print(bold(yellow_bg("average accuracy: ", avg_accuracy)))

        net_mat = net_mat.astype('int')
        print(bold(cyan(net_mat)))
        plot_confusion_matrix(net_mat,show_normed=True, colorbar=True, show_absolute=True, cmap='Blues')

        plt.show()

        roc_score=roc_auc_score(Y,Kn.predict_proba(X), multi_class='ovr')
        print(bold(yellow_bg("ROC_AUC_SCORE: ",roc_score)))

        print(bold("Classification Report:\n"))
        report=classification_report(y_ori, y_pre,output_dict=True)
        report_Df=pd.DataFrame(report)
        print(bold(blue(report_Df)))
        sns.heatmap(report_Df.T,annot=True)
        return avg_accuracy
avg_accuracy_knn=KNN(Independent_Attributes,Dependent_Attribute)
```